

SGD Resample

September 30, 2019

1 Discrete Algorithm

let $\mathbf{x} \in \mathbb{R}^d$ to denote the input feature and we wish to do regression task, i.e. to approximate response variable $y = g(x) \in \mathbb{R}$ using a three-layer neural network $\hat{y} = f(x)$. The architecture of the neural network is defined as follows:

$$\begin{aligned}\mathbf{h}_1 &= \sigma(\mathbf{W}^{(1)}\mathbf{x}) \\ \mathbf{h}_2 &= \sigma(\mathbf{W}^{(2)}\mathbf{h}_1) \\ \hat{y} &= \mathbf{W}^{(3)}\mathbf{h}_2\end{aligned}\tag{1}$$

, where $\mathbf{W}^{(i)} \in \mathbb{R}^{m_i \times m_{i-1}}$ is the weight matrix in i -th layer, and m_i is the hidden size in layer i , here the input layer is layer 0 and $m_0 = d$.

The SGD resample algorithm is doing as follows

1. minimize empirical loss w.r.t. last layer, the loss function is

$$\mathcal{L}(\mathbf{W}^{(3)}) = \frac{1}{n} \sum_{i=1}^n (f(x) - y)^2 + \frac{\lambda_3}{1} \|\mathbf{W}^{(3)}\|_F$$

2. resample weights defined in layer 2: minimize the regularization terms w.r.t. weights connected to layer 2 and fix function value in layer 3 invariant:

$$\mathcal{L}(\mathbf{W}^{(2)}, \mathbf{W}^{(3)}) = \frac{\gamma}{n} \sum_{i=1}^n (\hat{y} - \hat{y}^{old})^2 + \frac{\lambda_3}{1} \|\mathbf{W}^{(3)}\|_F + \frac{\lambda_2}{m_2} \|\mathbf{W}^{(2)}\|_F$$

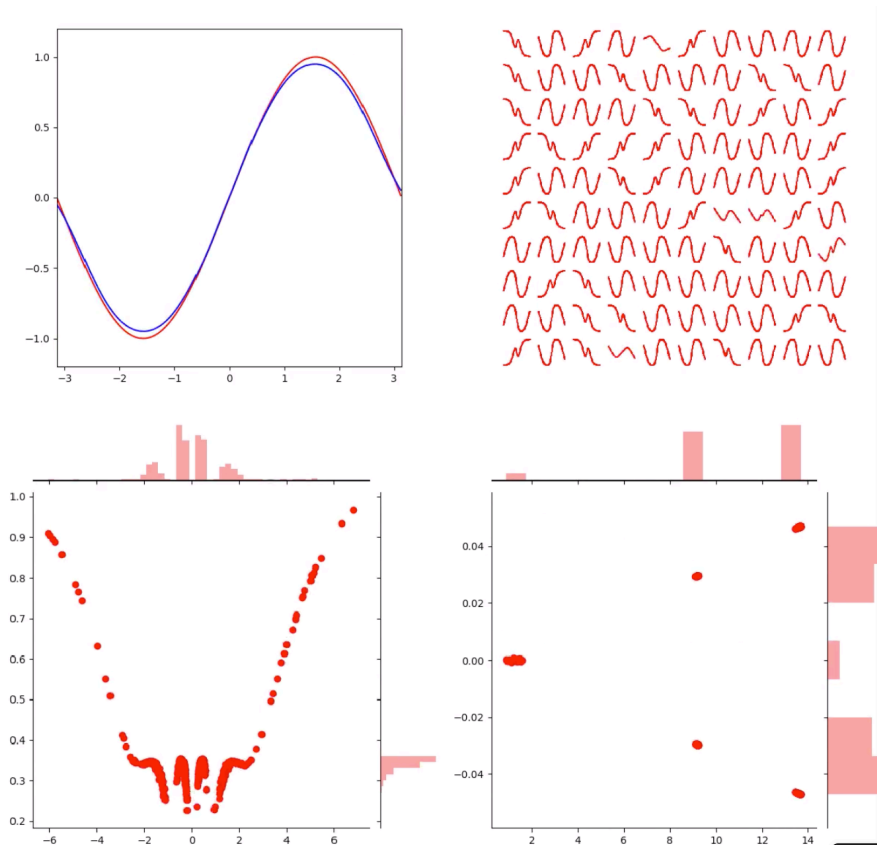
, where γ are hyperparameters that control the constraint of invariance of function value, x^{old} denote the respective value x in original neural network.

3. resample weights defined in layer 1: minimize the regularization terms w.r.t. weights connected to layer 1 and fix function value in layer 2 invariant:

$$\mathcal{L}(\mathbf{W}^{(1)}, \mathbf{W}^{(2)}) = \frac{\gamma}{n} \sum_{i=1}^n \|\mathbf{h}_2 - \mathbf{h}_2^{old}\|_2^2 + \frac{\lambda_2}{m_2} \|\mathbf{W}^{(2)}\|_F + \frac{\lambda_1}{m_1} \|\mathbf{W}^{(1)}\|_F$$

1.1 Joint Training Results

The stationary distributions of parameters learned by standard sgd are as follows:



| | |
|---|--|
| <p><i>prediction function</i> x: input feature, y: output red: groundtruth, blue: prediction</p> | <p><i>Samples of θ_2</i> x: input feature, y: response of the neuron θ_2 in layer 2 each subplot: a sample of θ_2</p> |
| <p><i>dist of θ_1 and $\omega(\theta_1)$</i> x: θ_1, y: $\omega(\theta_1) = \sqrt{\int \theta_2(\theta_1) ^2 \rho(d\theta_2)}$ top and right: marginal histogram</p> | <p><i>dist of $\ \theta_2\ _2$ and $w(\theta_2)$</i> x: $\ \theta_2\ _2$, y: $w(\theta_2)$ top and right: marginal histogram</p> |

final regression loss $0.3e-3$, regularization loss: $3.0e-3$.

1.2 Steps towards alternative training

1. use well trained $\mathbf{W}^{(1)}$, and fix them during training: only alternating between step [1] and step [2] in discrete algorithm (video fixedtheta1.mp4)
 - final regression loss: $0.3e-3$, regularization loss: $2.7e-3$
 - could share similar distribution w.r.t. the conventional SGD algorithm.
 - attained similar regression loss and smaller regularization loss (and regression loss is more stable)
2. use well trained $\mathbf{W}^{(1)}$ as initialization and alternating between step [1], [2], [3] (video welltheta1.mp4)
 - final regression loss: $1.4e-3$, regularization loss $2.4e-3$.
 - it seems that the distribution of θ_2 converges quickly but the regression loss converge very slowly.
3. use random $\mathbf{W}^{(1)}$ as initialization, do the original algorithm
 - final regression loss: $1e-2$
 - could not converge (both regression loss and distribution)